

Lecture Notes for **Machine Learning in Python**

Professor Eric Larson
Neural Network Optimization and Activation

Class Logistics and Agenda

- Agenda:
 - More optimization techniques
 - Momentum
 - Adaptive learning rates
 - Initialization
 - More activations: Tanh, ReLU, SiLU
 - Programming Examples

Class Overview, by topic

Table Data
Visualization

Numpy, Pandas, Seaborn
Overviews with some in-depth discussion

Dimension
Reduction and
Image Processing

Scikit-learn, Scikit Image,
Intuition only, Some mathematics

Linear and
Logistic
Regression

Numpy, Recreate API for Scikit-learn
Detailed mathematics for simple optimization
intuition for advanced optimization

Neural Networks
and Back Prop.

Numpy
Detailed mathematics for NN operations

Wide and Deep
Networks

Convolutional
Networks

Recurrent
Networks

Keras, Tensorflow
Intuition, Detailed implement.

Ethics in
Language Models

ConceptNet
Case studies

Mini-batching

- Numerous instances to find one gradient update
 - **solution:** mini-batch


← all data →

	batch 1	batch 2	batch 3	batch 4	batch 5	batch 6	batch 7	batch 8	batch 9
Epoch 1									
Epoch 2									
Epoch 3									
Epoch 4									
...									

shuffle ordering each epoch and update W 's after each batch

- **new problem:** mini-batch gradient updates erratic
 - **solutions:**
 - momentum
 - adaptive learning steps (cooling)

Momentum and Cooling Intuition




Momentum

$$\mathbf{W}_{k+1} = \mathbf{W}_k - \rho_k$$

- Momentum $\rho_k = \alpha \nabla J(\mathbf{W}_k) + \beta \nabla J(\mathbf{W}_{k-1})$



- Nesterov's Accelerated Gradient $\rho_k = \underbrace{\beta \nabla J(\mathbf{W}_k + \alpha \nabla J(\mathbf{W}_{k-1})) + \alpha \nabla J(\mathbf{W}_{k-1})}_{\text{step twice}}$



Adaptive Strategy: Cooling

- Fixed Reduction at Each Epoch
- Adjust on Plateau
 - make smaller if when J rapidly changes
 - make bigger when J not changing much

$$\eta_e = \eta_0^{(1+e \cdot \epsilon)}$$



07. MLP Neural Networks.ipynb

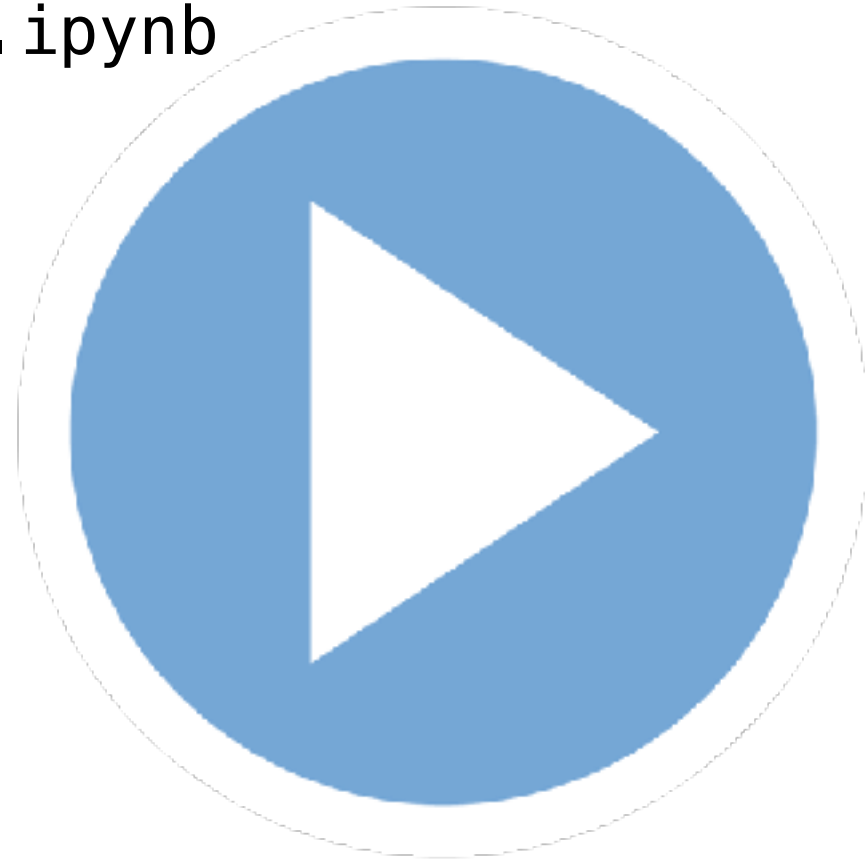
comparison:

mini-batch

momentum

adaptive learning

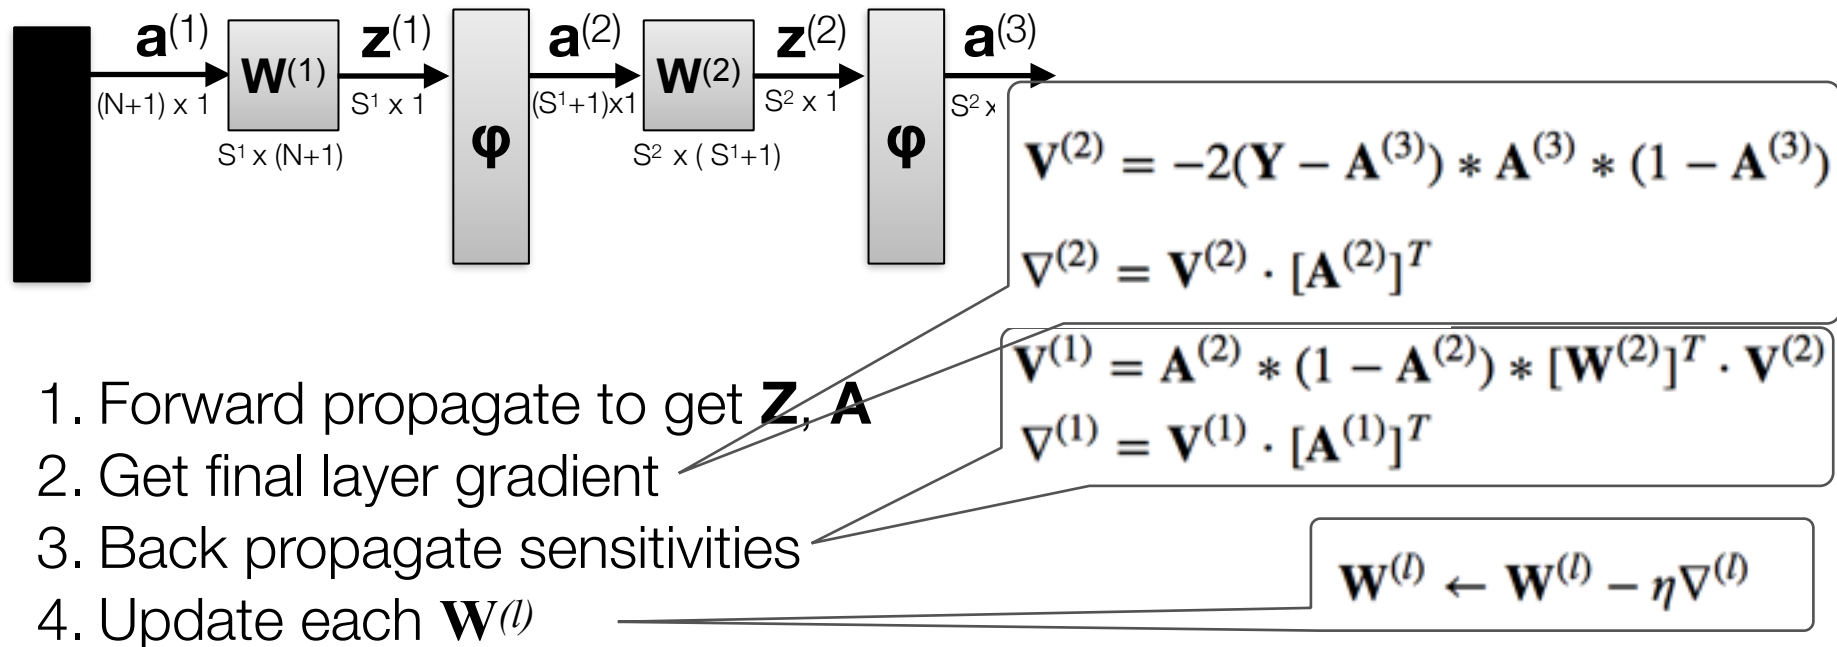
L-BFGS



Objective Function



Changing the Objective Function



• Self Test:

True or False: If we change the cost function, $J(\mathbf{W})$, we only need to update the final layer sensitivity calculation, $\mathbf{V}^{(2)}$, of the back propagation steps. The remainder of the algorithm is unchanged.


- A. True
- B. False

Practical Implementation of Architectures

- MSE

$$J(\mathbf{W}) = \sum_k^M (\mathbf{y}^{(k)} - \mathbf{a}^{(L)})^2$$

least squares objective,
tends to slow training initially




Run

Practical Implementation of Architectures

- MSE

$$J(\mathbf{W}) = \sum_k^M (\mathbf{y}^{(k)} - \mathbf{a}^{(L)})^2$$

least squares objective,
tends to slow training initially




Run

Practical Implementation of Architectures

- Negative of MLE: **Binary Cross entropy**

$$J(\mathbf{W}) = - [\mathbf{y}^{(i)} \ln([\mathbf{a}^{(L+1)}]^{(i)}) + (1 - \mathbf{y}^{(i)}) \ln(1 - [\mathbf{a}^{(L+1)}]^{(i)})]$$

speeds up initial training




Neural Networks and Deep Learning, Michael Nielson, 2015

Practical Implementation of Architectures

- Negative of MLE: **Binary Cross entropy**

$$J(\mathbf{W}) = - [\mathbf{y}^{(i)} \ln([\mathbf{a}^{(L+1)}]^{(i)}) + (1 - \mathbf{y}^{(i)}) \ln(1 - [\mathbf{a}^{(L+1)}]^{(i)})]$$

speeds up initial training



Practical Implementation of Architectures

$$J(\mathbf{W}) = - [\mathbf{y}^{(i)} \ln([\mathbf{a}^{(L+1)}]^{(i)}) + (1 - \mathbf{y}^{(i)}) \ln(1 - [\mathbf{a}^{(L+1)}]^{(i)})] \quad \text{speeds up initial training}$$

$$\begin{aligned} \left[\frac{\partial J(\mathbf{W})}{\partial \mathbf{z}^{(L)}} \right]^{(i)} &= - \frac{\partial}{\partial \mathbf{z}^{(L)}} [\mathbf{y}^{(i)} \ln([\mathbf{a}^{(L+1)}]^{(i)}) + (1 - \mathbf{y}^{(i)}) \ln(1 - [\mathbf{a}^{(L+1)}]^{(i)})] \\ &= - \left[\mathbf{y}^{(i)} \frac{\partial}{\partial \mathbf{z}^{(L)}} (\ln([\mathbf{a}^{(L+1)}]^{(i)})) + (1 - \mathbf{y}^{(i)}) \frac{\partial}{\partial \mathbf{z}^{(L)}} (\ln(1 - [\mathbf{a}^{(L+1)}]^{(i)})) \right] \\ &= - \left[\mathbf{y}^{(i)} \frac{1}{[\mathbf{a}^{(L+1)}]^{(i)}} ([\mathbf{a}^{(L+1)}]^{(i)} (1 - [\mathbf{a}^{(L+1)}]^{(i)})) + \frac{(1 - \mathbf{y}^{(i)})}{1 - [\mathbf{a}^{(L+1)}]^{(i)}} \left(- \frac{\partial}{\partial \mathbf{z}^{(L)}} [\mathbf{a}^{(L+1)}]^{(i)} \right) \right] \\ &= - \left[\mathbf{y}^{(i)} (1 - [\mathbf{a}^{(L+1)}]^{(i)}) - \frac{(1 - \mathbf{y}^{(i)})}{1 - [\mathbf{a}^{(L+1)}]^{(i)}} ([\mathbf{a}^{(L+1)}]^{(i)} (1 - [\mathbf{a}^{(L+1)}]^{(i)})) \right] \\ &= - [\mathbf{y}^{(i)} (1 - [\mathbf{a}^{(L+1)}]^{(i)}) - (1 - \mathbf{y}^{(i)}) ([\mathbf{a}^{(L+1)}]^{(i)})] \\ &= - [\mathbf{y}^{(i)} - \mathbf{y}^{(i)} [\mathbf{a}^{(L+1)}]^{(i)} - [\mathbf{a}^{(L+1)}]^{(i)} + [\mathbf{a}^{(L+1)}]^{(i)} \mathbf{y}^{(i)}] = [\mathbf{a}^{(L+1)}]^{(i)} - \mathbf{y}^{(i)} \end{aligned}$$

$$\mathbf{V}^{(2)} = -2(\mathbf{Y} - \mathbf{A}^{(3)}) * \mathbf{A}^{(3)} * (1 - \mathbf{A}^{(3)}) \quad \text{old update}$$

Practical Implementation of Architectures

- Back to our old friend: **Cross entropy**

$$J(\mathbf{W}) = - \left[\mathbf{y}^{(i)} \ln([\mathbf{a}^{(L+1)}]^{(i)}) + (1 - \mathbf{y}^{(i)}) \ln(1 - [\mathbf{a}^{(L+1)}]^{(i)}) \right]$$

speeds up
initial training

$$\left[\frac{\partial J(\mathbf{W})}{\partial \mathbf{z}^{(L)}} \right]^{(i)} = ([\mathbf{a}^{(L+1)}]^{(i)} - \mathbf{y}^{(i)})$$

$$\left[\frac{\partial J(\mathbf{W})}{\partial \mathbf{z}^{(2)}} \right]^{(i)} = ([\mathbf{a}^{(3)}]^{(i)} - \mathbf{y}^{(i)})$$

$$\mathbf{V}^{(2)} = \mathbf{A}^{(3)} - \mathbf{Y}$$

new update

```
# vectorized backpropagation
V2 = (A3 - Y_enc) # <- this is only line t
V1 = A2 * (1 - A2) * (W2.T @ V2)

grad2 = V2 @ A2.T
grad1 = V1[1:,:] @ A1.T
```

$$\mathbf{V}^{(2)} = -2(\mathbf{Y} - \mathbf{A}^{(3)}) * \mathbf{A}^{(3)} * (1 - \mathbf{A}^{(3)})$$

old update

bp-5

cross entropy



Practical Implementation of Architectures

SQL programmers be like




Formative Self Test

- for adding Gaussian distributions, variances add together

$\mathbf{a}^{(L+1)} = \phi(\mathbf{W}^{(L)}\mathbf{a}^{(L)})$ assume each element of \mathbf{a} is Gaussian

- If you initialized the weights, \mathbf{W} , with too large variance, you would expect the output of the neuron, $\mathbf{a}^{(L+1)}$, to be:
 - A. saturated to “1”
 - B. saturated to “0”
 - C. could either be saturated to “0” or “1”
 - D. would not be saturated




Formative Self Test

- for adding Gaussian distributions, variances add together

$\mathbf{a}^{(L+1)} = \phi(\mathbf{W}^{(L)}\mathbf{a}^{(L)})$ assume each element of \mathbf{a} is Gaussian

- What is the derivative of a saturated sigmoid neuron?
 - A. zero
 - B. one
 - C. $a * (1-a)$
 - D. it depends



Practical Implementation of Architectures

- **Weight initialization**


- try not to **saturate** your neurons right away!

$$\mathbf{a}^{(L+1)} = \phi(\mathbf{z}^{(L)})$$

$$\mathbf{z}^{(L)} = \mathbf{W}^{(L)} \mathbf{a}^{(L)}$$



each row is summed before sigmoid



want each $z^{(L)}$ to be between $-\epsilon < \Sigma < \epsilon$ for no saturation

solution: squash initial weights magnitude

- one choice: each element of \mathbf{W} selected from a Gaussian with **zero mean** and **specific standard deviation**

$$w_{ij}^{(L)} \leftarrow \mathcal{N}\left(0, \sqrt{\frac{1}{n^{(L)}}}\right)$$

For a sigmoid, want $-\epsilon < z_i^{(L)} < \epsilon$

$\epsilon=4$

More Weight Initialization

Understanding the difficulty of training deep feedforward neural networks

Xavier Glorot JMLR 2010 Yoshua Bengio
DIRO, Université de Montréal, Montréal, Québec, Canada

Goal: We should not saturate **feedforward** or **back propagated** variance

Relate variance of current layer to variance in z , so $\sigma(z_i^{(L)})$ isn't saturated

try not to saturate z $z_i^{(L)} = \sum_j^{n^{(L)}} w_{ij} a_j^{(L)}$ *break down feed forward by each multiply*

$$\text{Var}[z_i^{(L)}] = \sum_j^{n^{(L)}} \underbrace{E[w_{ij}]^2 \text{Var}[a_j^{(L)}] + \text{Var}[w_{ij}] E[a_j^{(L)}]^2}_{0, \text{ if uncorrelated}} + \underbrace{\text{Var}[w_{ij}] \text{Var}[a_j^{(L)}]}_{\approx 1}$$

Want to keep $\text{Var}[\cdot] \sim 1$ *assume i.i.d. expand variance calc*

$$\text{Var}[z_i^{(L)}] = 4 = n^{(L)} \text{Var}[w_{ij}] \text{Var}[a_j^{(L)}]$$

$$w_{ij}^{(L)} \approx \mathcal{N}\left(0, 4 \cdot \sqrt{\frac{1}{n^{(L)}}}\right)$$

forward
from data

More Weight Initialization

$$\text{Var}[z_i^{(L)}] = 4 = n^{(L)} \text{Var}[w_{ij}] \text{Var}[a_j^{(L)}]$$

$$w_{ij}^{(L)} \approx \mathcal{N} \left(0, 4 \cdot \sqrt{\frac{1}{n^{(L)}}} \right)$$

forward
from data

$$\mathbf{v}^{(L)} = \mathbf{a}^{(L)}(1 - \mathbf{a}^{(L)})\mathbf{W}^{(L)} \cdot \mathbf{v}^{(L+1)}$$

Similar for back prop.

$$\text{Var}[v_i^{(L)}] = n^{(L+1)} \text{Var}[w_{ij}] \text{Var}[v_j^{(L+1)} \cdot a_j^{(L)}(1 - a_j^{(L)})]$$

$$w_{ij}^{(L)} \approx \mathcal{N} \left(0, 4 \cdot \sqrt{\frac{1}{n^{(L+1)}}} \right)$$

backward
from sensitivity

$$w_{ij}^{(L)} \approx \mathcal{N} \left(0, 4 \cdot \sqrt{\frac{2}{n^{(L)} + n^{(L+1)}}} \right)$$

compromise


More Weight Initialization

Understanding the difficulty of training deep feedforward neural networks

Xavier Glorot

DIRO, Université de Montréal, Montréal, Québec, Canada

Yoshua Bengio




Starting gradient histograms
per layer
standard initialization

Starting gradient histograms
per layer
Glorot initialization


Figure 7: *Back-propagated gradients normalized histograms with hyperbolic tangent activation, with standard (top) vs normalized (bottom) initialization. Top: 0-peak decreases for higher layers.*

Smarter Weight Initialization



New Activation: Hyperbolic Tangent

- Basically a sigmoid from -1 to 1





$$\phi(z) = \frac{\sinh(z)}{\cosh(z)} = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

$$\frac{\partial}{\partial z}\phi(z) = \text{sech}^2(z)$$

New Activation: ReLU

- A new nonlinearity: **rectified linear units**




$$\phi(z) = \begin{cases} z, & \text{if } z > 0 \\ 0, & \text{else} \end{cases}$$

it has the advantage of **large gradients** and **extremely simple** derivative

$$\nabla \phi(z) = \begin{cases} 1, & \text{if } z > 0 \\ 0, & \text{else} \end{cases}$$

Other Activation Functions



- Sigmoid Weighted Linear Unit **SiLU**
 - also called Swish
- Mixing of sigmoid, σ , and ReLU

$$\phi(z) = z \cdot \sigma(z)$$

$$\frac{\partial \phi(z)}{\partial z} = \phi(z) + \sigma(z)[1 - \phi(z)]$$

$$= a^{(l+1)} + \sigma(z^{(l)}) \cdot [1 - a^{(l+1)}]$$




Figure 1: The Swish activation function.

Ramachandran P, Zoph B, Le QV. Swish: a Self-Gated Activation Function. arXiv preprint arXiv:1710.05941. 2017 Oct 16

Elfwing, Stefan, Eiji Uchibe, and Kenji Doya. "Sigmoid-weighted linear units for neural network function approximation in reinforcement learning." Neural Networks (2018).

Derivative Calculation:

$$= \sigma(x) + x \cdot \sigma(x)(1 - \sigma(x))$$

$$= \sigma(x) + x \cdot \sigma(x) - x \cdot \sigma(x)^2$$

$$= x \cdot \sigma(x) + \sigma(x)(1 - x \cdot \sigma(x))$$

Glorot and He Initialization

We have solved this assuming the activation output is in the range -4 to 4 (for a sigmoid) and assuming that x is distributed Gaussian

This range, epsilon, is different depending on the activation and assuming Gaussian or Uniform

Uniform

Gaussian

Tanh

$$w_{ij}^{(L)} = \sqrt{\frac{6}{n^{(L)} + n^{(L+1)}}}$$

$$w_{ij}^{(L)} = \sqrt{\frac{2}{n^{(L)} + n^{(L+1)}}}$$

Sigmoid

$$w_{ij}^{(L)} = 4\sqrt{\frac{6}{n^{(L)} + n^{(L+1)}}}$$

$$w_{ij}^{(L)} = 4\sqrt{\frac{2}{n^{(L)} + n^{(L+1)}}}$$

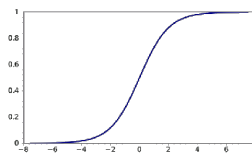
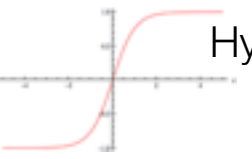
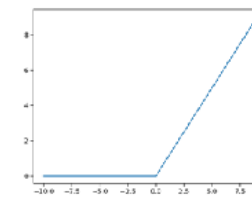
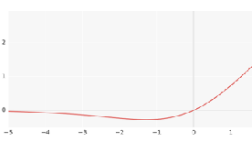
ReLU

$$w_{ij}^{(L)} = \sqrt{2}\sqrt{\frac{6}{n^{(L)} + n^{(L+1)}}}$$

$$w_{ij}^{(L)} = \sqrt{2}\sqrt{\frac{2}{n^{(L)} + n^{(L+1)}}}$$

Summarized by Glorot and He

Activations Summary

	Definition	Derivative	Weight Init (Uniform Bounds)
 <p>Sigmoid</p>	$\phi(z) = \frac{1}{1 + e^{-z}}$	$\nabla \phi(z) = a(1 - a)$	$w_{ij}^{(L)} = 4\sqrt{\frac{6}{n^{(L)} + n^{(L+1)}}}$
 <p>Hyperbolic Tangent</p>	$\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	$\nabla \phi(z) = \frac{4}{(e^z + e^{-z})^2}$	$w_{ij}^{(L)} = \sqrt{\frac{6}{n^{(L)} + n^{(L+1)}}}$
 <p>ReLU</p>	$\phi(z) = \begin{cases} z, & \text{if } z > 0 \\ 0, & \text{else} \end{cases}$	$\nabla \phi(z) = \begin{cases} 1, & \text{if } z > 0 \\ 0, & \text{else} \end{cases}$	$w_{ij}^{(L)} = \sqrt{2}\sqrt{\frac{6}{n^{(L)} + n^{(L+1)}}}$
 <p>SiLU</p>	$\phi(z) = \frac{z}{1 + e^{-z}}$	$\nabla \phi(z) = a + \frac{(1 - a)}{1 + e^{-z}}$	$w_{ij}^{(L)} = \sqrt{2}\sqrt{\frac{6}{n^{(L)} + n^{(L+1)}}}$

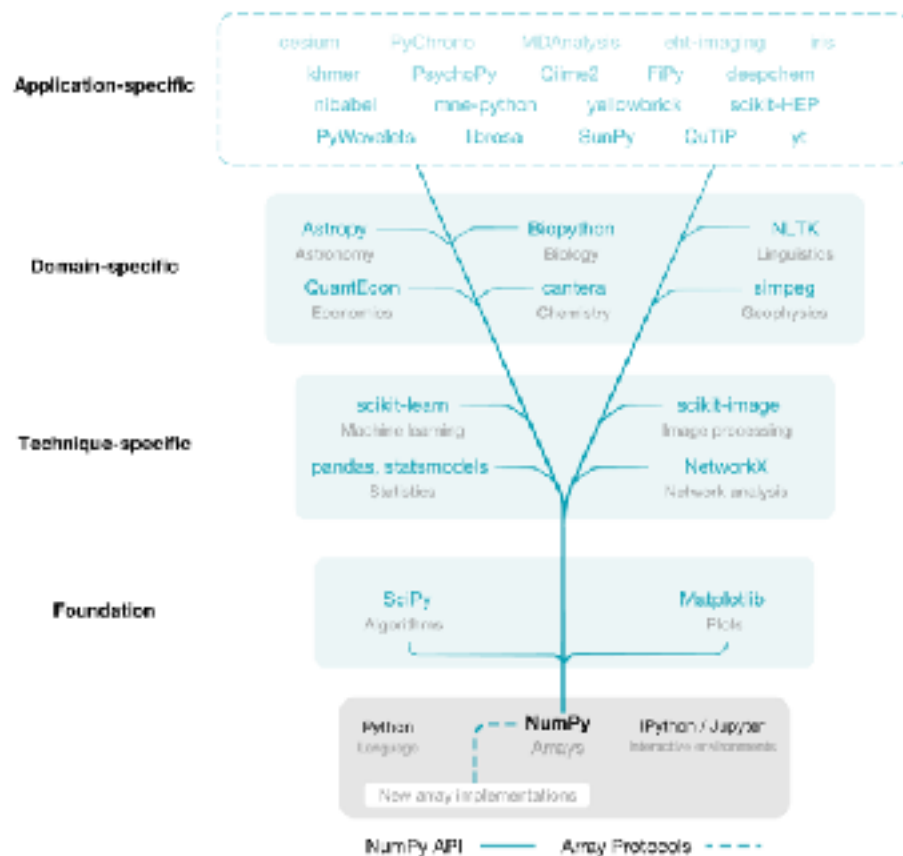


ReLU Nonlinearities
Important for deep networks

Fig. 2: NumPy is the base of the scientific Python ecosystem.


From: Array programming with NumPy

More Adaptive Optimization



Be adaptive based on Gradient Magnitude?

- Decelerate down regions that are steep
- Accelerate on plateaus




Use **inverse** of
magnitude of **gradient**!

Also **accumulate inverse** to be robust to
abrupt changes in **steepness**... momentum!!


Be adaptive based on Gradient Magnitude?

Inverse magnitude of gradient in multiple directions?

$$\mathbf{W}_{k+1} \leftarrow \mathbf{W}_k + \eta \frac{1}{\sqrt{\mathbf{G}_k}} \odot \nabla J(\mathbf{W}_k)$$



$$\mathbf{G}_k = \nabla J(\mathbf{W}_k) \odot \nabla J(\mathbf{W}_k)$$



Common Adaptive Strategies $\mathbf{W}_{k+1} = \mathbf{W}_k - \eta \cdot \rho_k$

Adjust each element of gradient by the steepness

- AdaGrad
all operations are per element
$$\rho_k = \frac{1}{\sqrt{\mathbf{G}_k + \epsilon}} \odot \nabla J(\mathbf{W}_k) \quad \text{where} \quad \mathbf{G}_k = \mathbf{G}_{k-1} + \nabla J(\mathbf{W}_k) \odot \nabla J(\mathbf{W}_k)$$
- RMSProp
all operations are per element
$$\rho_k = \frac{1}{\sqrt{V_k + \epsilon}} \odot \nabla J(\mathbf{W}_k) \quad \begin{aligned} \mathbf{G}_k &= \nabla J(\mathbf{W}_k) \odot \nabla J(\mathbf{W}_k) \\ \mathbf{V}_k &= \gamma \cdot \mathbf{V}_{k-1} + (1 - \gamma) \cdot \mathbf{G}_k \end{aligned}$$
- AdaDelta
all operations are per element
$$\rho_k = \frac{\mathbf{M}_k}{\sqrt{\mathbf{V}_k + \epsilon}} \quad \mathbf{M}_k = \gamma \cdot \mathbf{M}_k + (1 - \gamma) \cdot \nabla J(\mathbf{W}_k)$$
- AdaM
 \mathbf{G} updates with decaying momentum of J and J^2
- NAdaM
same as Adam, but with nesterov's acceleration

None of these are “**one-size-fits-all**” because the space of neural network **optimization varies** by problem, ADAM is **popular** but **not a panacea**