

Lecture Notes for **Machine Learning in Python**

Professor Eric Larson
Basic Convolutional Neural Networks

Logistics and Agenda

- Logistics
 - Wide/Deep due soon!
 - Remember: Feel free to turn in late for partial credit.
- Agenda
 - Wide/Deep Finish Demo and Town Hall
 - Basic CNN architectures and Demo

Class Overview, by topic

Table Data
Visualization

Numpy, Pandas, Seaborn
Overviews with some in-depth discussion

Dimension
Reduction and
Image Processing

Scikit-learn, Scikit Image,
Intuition only, Some mathematics

Linear and
Logistic
Regression

Numpy, Recreate API for Scikit-learn
Detailed mathematics for simple optimization
intuition for advanced optimization

Neural Networks
and Back Prop.

Numpy
Detailed mathematics for NN operations

Wide and Deep
Networks

Convolutional
Networks

Recurrent
Networks


Keras, Tensorflow
Intuition, Detailed implement.

Ethics in
Language Models

ConceptNet
Case studies

Last Time:

- Deep refers to increasingly smaller hidden layers
- Embed into sparse representations via ReLU

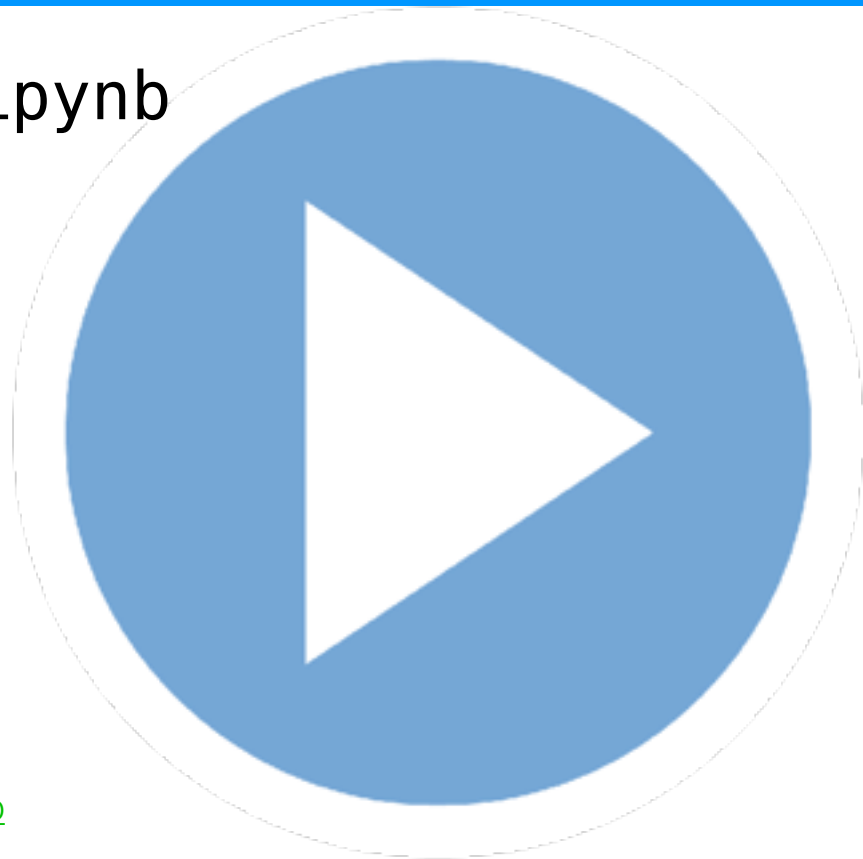


10. Keras Wide and Deep.ipynb

The awful dataset:
Toy Census Data Example


Other tutorials:

https://www.tensorflow.org/tutorials/wide_and_deep



Town Hall, Wide and Deep Networks

When $p < 0.05$



Convolutional Neural Networks



IN CS, IT CAN BE HARD TO EXPLAIN
THE DIFFERENCE BETWEEN THE EASY
AND THE VIRTUALLY IMPOSSIBLE.

Reminder: Convolution

$$\sum \left(\mathbf{I} \left[i \pm \frac{r}{2}, j \pm \frac{c}{2} \right] \odot \mathbf{k} \right) = \mathbf{O}[i, j] \quad \begin{array}{l} \text{output image} \\ \text{at pixel } i, j \end{array}$$

input image at $r \times c$ range of
pixels centered in i, j

kernel of size, $r \times c$
usually $r=c$

0	0	0	0	0	0	0	0	0
0	1	2	3	4	12	9	8	0
0	5	2	3	4	12	9	8	0
0	5	2	1	4	10	9	8	0
0	7	2	1	4	12	7	8	0
0	7	2	1	4	14	9	8	0
0	5	2	3	4	12	7	8	0
0	5	2	1	4	12	9	8	0
0	0	0	0	0	0	0	0	0

input image, \mathbf{I}

1	2	1
2	4	2
1	2	1

kernel
filter, \mathbf{k}
3x3

20	21	36
...
...
...
...
...
...
...

output image, \mathbf{O}

Reminder: Convolution

$$\begin{array}{|c|c|} \hline O_{11} & O_{12} \\ \hline O_{21} & O_{22} \\ \hline \end{array} = \text{Convolution} \left(\begin{array}{|c|c|c|} \hline X_{11} & X_{12} & X_{13} \\ \hline X_{21} & X_{22} & X_{23} \\ \hline X_{31} & X_{32} & X_{33} \\ \hline \end{array}, \begin{array}{|c|c|} \hline F_{11} & F_{12} \\ \hline F_{21} & F_{22} \\ \hline \end{array} \right)$$

Output \odot
Input \mathbf{X}
Filter \mathbf{F}

X_{11}	X_{12}	X_{13}
X_{21}	X_{22}	X_{23}
X_{31}	X_{32}	X_{33}

Input \mathbf{X}

\otimes

F_{11}	F_{12}
F_{21}	F_{22}

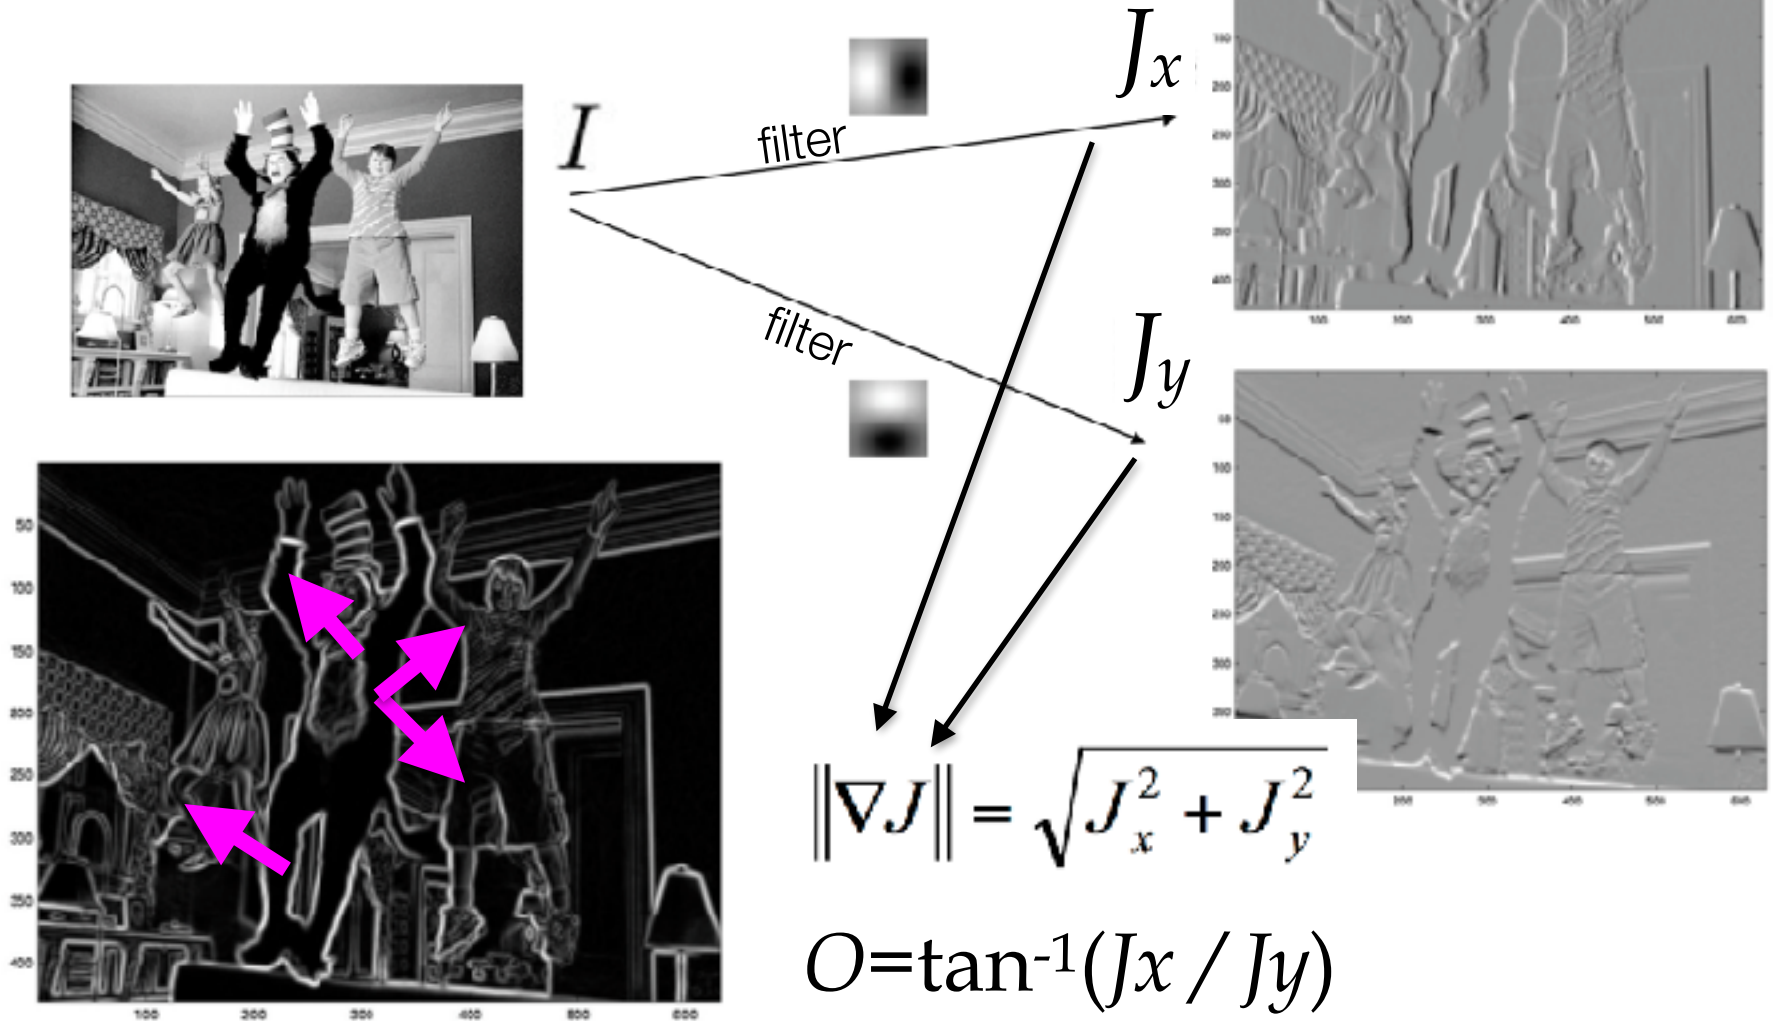
Filter \mathbf{F}

$X_{11}F_{11}$	$X_{12}F_{12}$	X_{13}
$X_{21}F_{21}$	$X_{22}F_{22}$	X_{23}
X_{31}	X_{32}	X_{33}

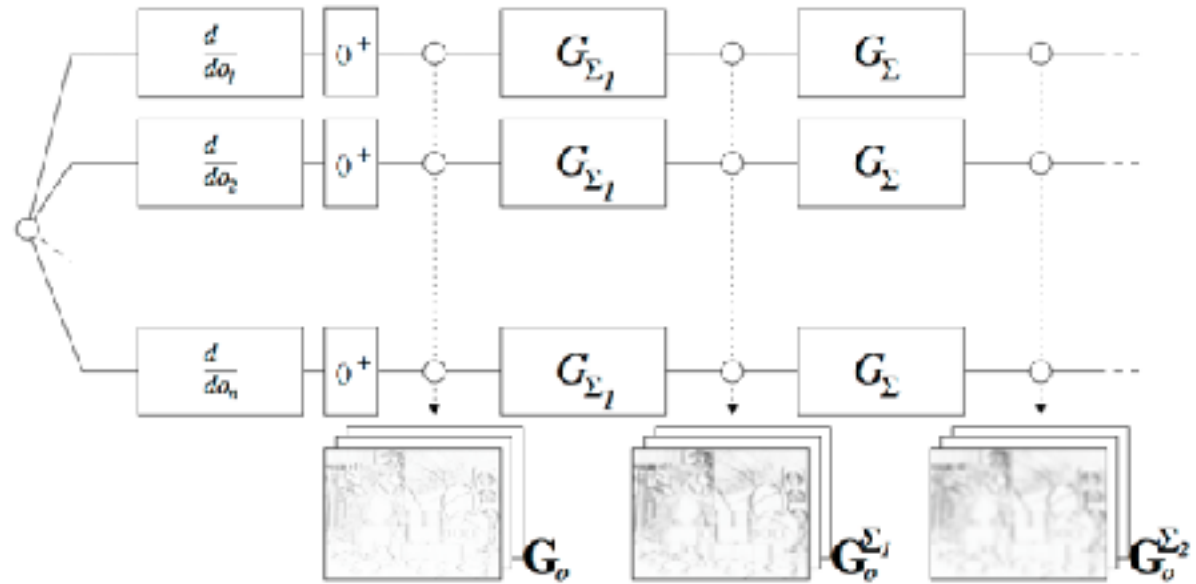
$$O_{11} = X_{11}F_{11} + X_{12}F_{12} + X_{21}F_{21} + X_{22}F_{22}$$

What we did before

- the gradient (2D derivative)



What we did before




take normalized histogram at point u, v

$$\tilde{\mathbf{h}}_{\Sigma}(u, v) = \left\| \left[\mathbf{G}_1^{\Sigma}(u, v), \dots, \mathbf{G}_H^{\Sigma}(u, v) \right]^{\top} \right\|$$

$$\mathcal{D}(u_0, v_0) = \begin{bmatrix} \tilde{\mathbf{h}}_{\Sigma_1}^{\top}(u_0, v_0), \\ \tilde{\mathbf{h}}_{\Sigma_1}^{\top}(\mathbf{l}_1(u_0, v_0, R_1)), \dots, \tilde{\mathbf{h}}_{\Sigma_1}^{\top}(\mathbf{l}_T(u_0, v_0, R_1)), \\ \tilde{\mathbf{h}}_{\Sigma_2}^{\top}(\mathbf{l}_1(u_0, v_0, R_2)), \dots, \tilde{\mathbf{h}}_{\Sigma_2}^{\top}(\mathbf{l}_T(u_0, v_0, R_2)), \end{bmatrix}$$


Tola et al. "Daisy: An efficient dense descriptor applied to wide-baseline stereo." Pattern Analysis and Machine Intelligence, IEEE Transactions

Anatomy of a convolutional network



CNN Overview

- First layer(s):
 - convolution
 - activation
 - pooling
 - Each pooling layer *can* make the input image “smaller”
 - allows for “Information Distillation”
 - less dependence on exact pixels
- Final layers are densely connected
 - typically multi-layer perceptrons




CNN Overview: Self Test

- First layer(s):
 - convolution
 - nonlinearity
 - pooling
 - Each pooling layer *can* make the input image “smaller”
 - allows for “Information Distillation”
 - less dependence on exact pixels
- Final layers are densely connected
 - typically multi-layer perceptrons
- Where are unstable gradients **most** problematic?
 - (A) During Convolution Layer(s) updates
 - (B) During Fully Connected Layer(s) updates
 - (C) Both A and B
 - (D) They are not a problem


CNN Filtering

- Why perform lots of filtering?
 - “recall” gabor filtering?



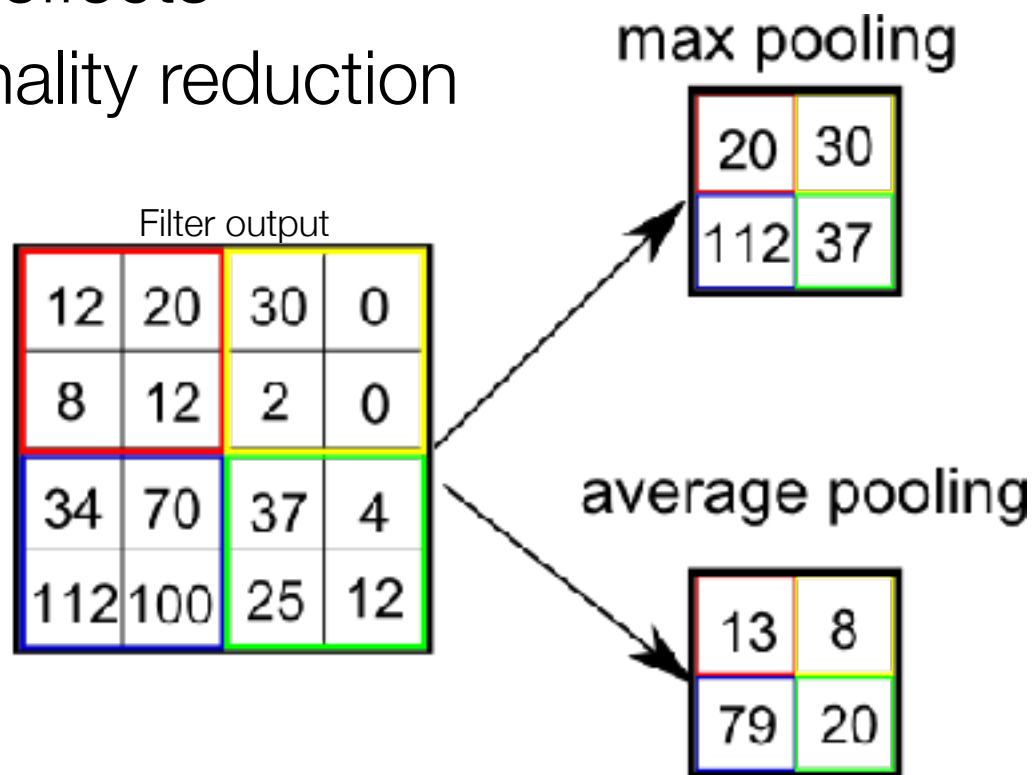


V1	Motion
V2	Stereo
V3	Color
V3a	Texture segregation
V3b	Segmentation, grouping
V4	Recognition
V7	Face recognition
MT	Attention
MST	Working memory/mental imagery

Convolution in a CNN

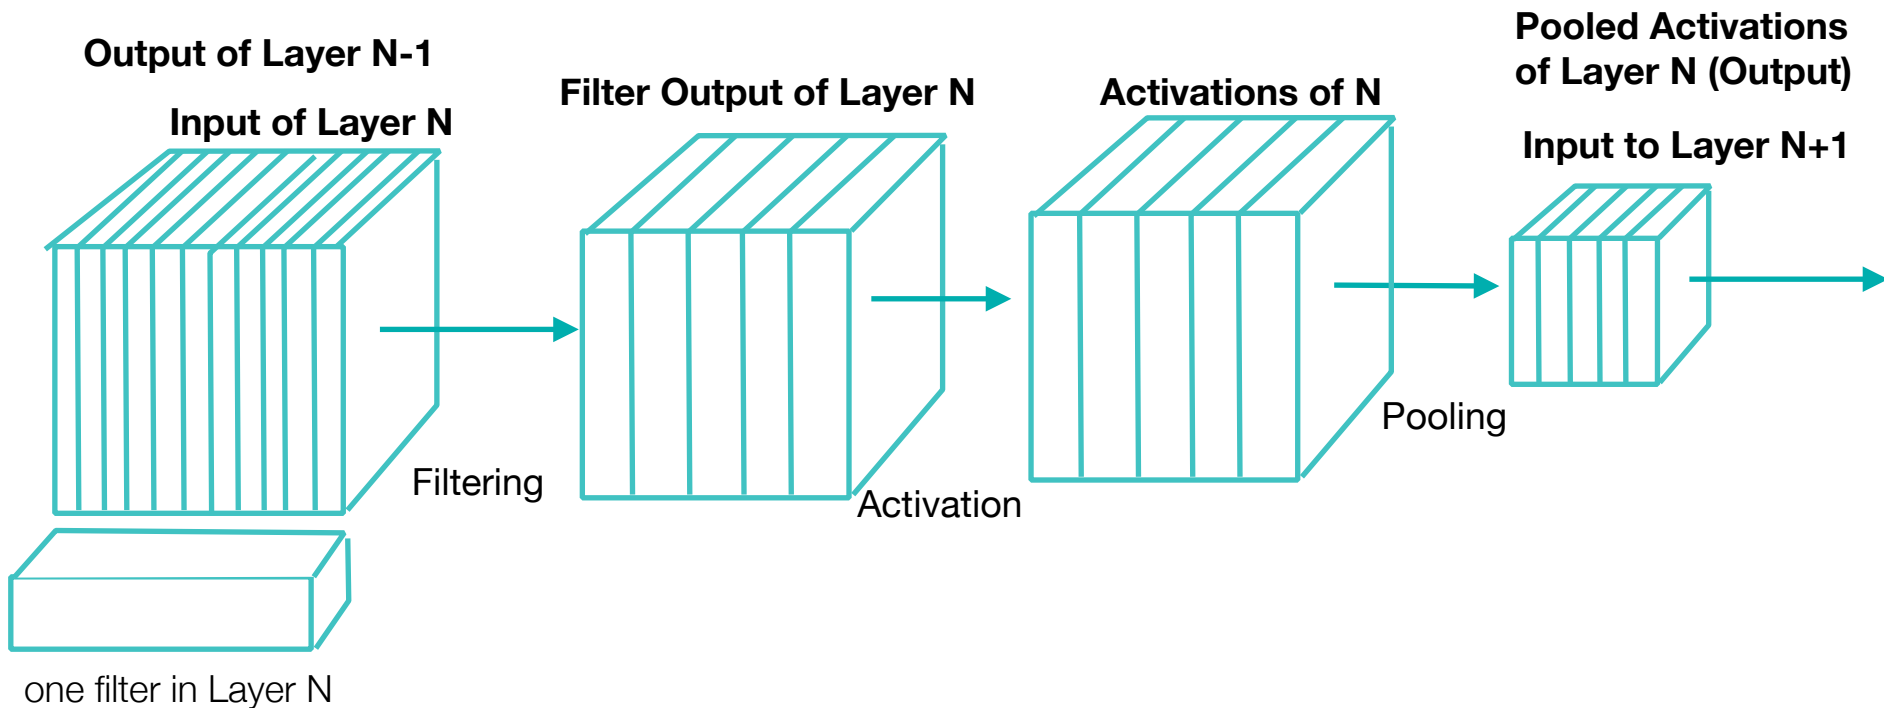


CNN Pooling

- Why perform pooling?
- Why max pooling?
 - reduce translation effects
 - **mostly**: dimensionality reduction

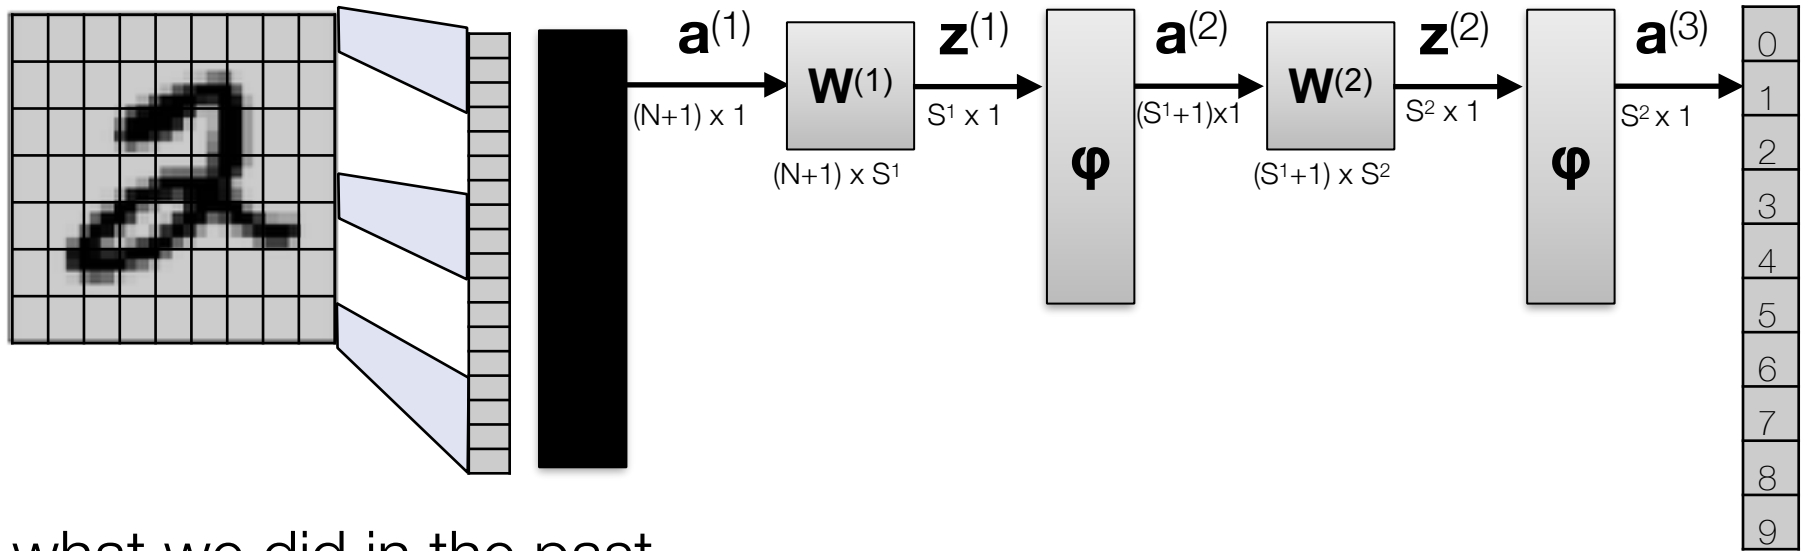


CNNs: Putting it together



Structure of Each Tensor: Channels x Rows x Columns

Simple Example: From Fully Connected to CNN




what we did in the past

If image is 9x9, and each fully connected layer is 20 hidden neurons wide, how many parameters are in this NN (ignore bias)?

$$(K^2 \times 20) + (20 \times 10) = 200 + 20 K^2$$


$$\text{for } 9 \times 9 = 200 + 20 \times 9^2 = 1,820 \text{ parameters}$$

Simple Example: From Fully Connected to CNN



If image is 9x9, and each fully connected layer is 20 hidden neurons wide, how many parameters are in this NN (ignore bias)?


3x3x3 = 27 weights



$$27 + (27 \times 20) + (20 \times 10) = 767$$

0
1
2
3
4
5
6
7
8
9

Simple Example: From Fully Connected to CNN




convolutional params

$N_k \times k^2$ ← filter dimension
 num filters


Input to MLP

$N_k \times (K^2/k^2)$

image dimension



CNN gradient



Derivative of max pool is easy:

for each input X_i


$\text{pool}'(X_i) = 1$ if X_i is max
0 else

$$X^{(L)} = \text{pool}(\sigma(O))$$

Derivative of convolution is more involved:

$$\begin{pmatrix} O_{11} & O_{12} \\ O_{21} & O_{22} \end{pmatrix} = \text{Convolution} \left(\begin{pmatrix} X_{11} & X_{12} & X_{13} \\ X_{21} & X_{22} & X_{23} \\ X_{31} & X_{32} & X_{33} \end{pmatrix}, \begin{pmatrix} F_{11} & F_{12} \\ F_{21} & F_{22} \end{pmatrix} \right)$$

Output O Input X Filter F



Next Lecture

- More CNN architectures and CNN history